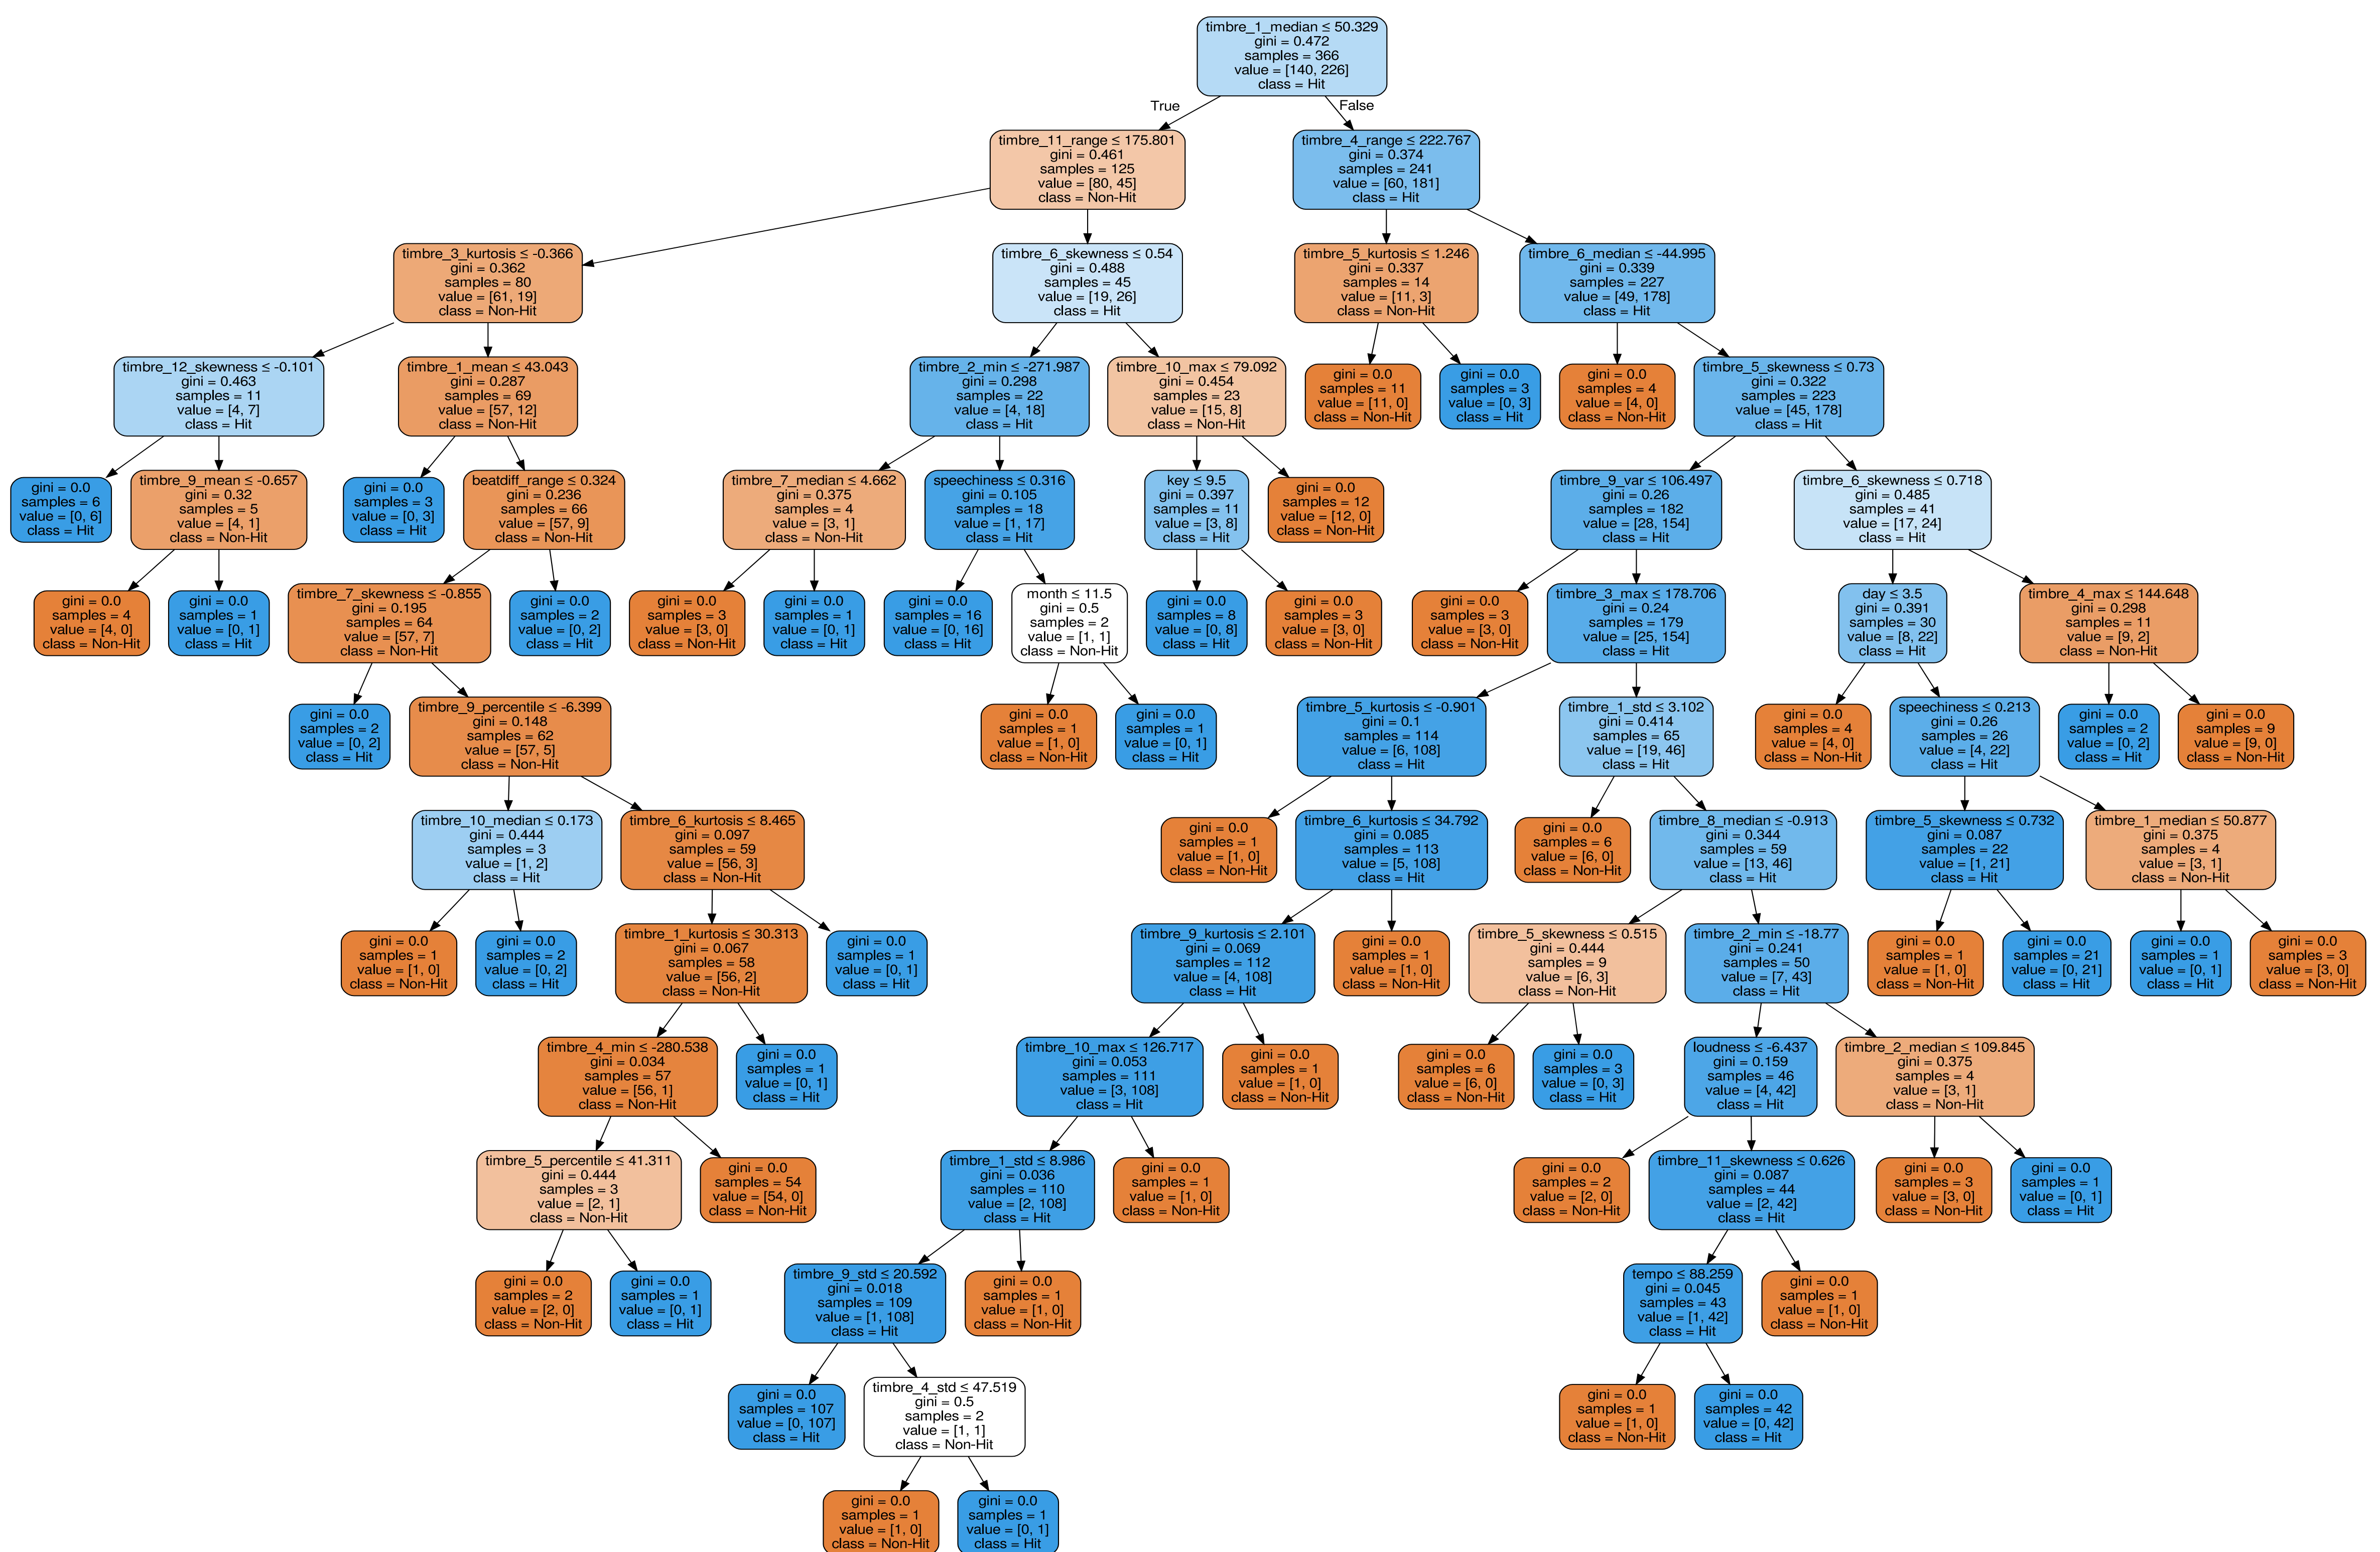


Investigating Hit Song Prediction

CARL DATH, CARL NYSTRÖMER, JAKOB VYTH & JOEL EKELOF

Supervised by Bob Sturm

- We are replicating the work presented in Herremans et al. “Dance Hit Song Prediction”, J. New Music Research 2014.
- Herremans et al. use machine learning to predict dance hit songs using features extracted from music recordings.
- They claim that their models “*can successfully predict if a dance song is going to be a top 10 hit versus a lower positioned dance song*”.
- We want to know what their models are doing and how they can be fooled.
- We have replicated their experiments.
- We are going to try to break their model by manipulating noise and get the model to classify it as a hit.



Decision tree trained on all features

Table 1: Results with 10-fold validation (accuracy) from Herremans et al.

Accuracy (%)	D1		D2		D3	
	-	FS	-	FS	-	FS
C4.5	<i>57.05</i>	<i>58.25</i>	<i>54.95</i>	<i>54.67</i>	<i>54.58</i>	<i>54.74</i>
RIPPER	<i>60.95</i>	<i>62.43</i>	<i>56.69</i>	<i>56.42</i>	<i>57.18</i>	<i>56.41</i>
Naive Bayes	65	65	<i>60.22</i>	<i>58.78</i>	<i>59.57</i>	<i>59.18</i>
Logistic regression	64.65	64	62.64	60.6	60.12	59.75
SVM (Polynomial)	64.97	64.7	61.55	61.6	61.04	61.07
SVM (RBF)	64.7	64.63	<i>59.8</i>	<i>59.89</i>	60.8	60.76

FS = feature selection, $p < 0.01$: italic, $p > 0.05$: bold, best: bold.